

# Mellors MSDS610 Week6 Assignment - Notebook 1: Splitting and Running the Model

## Loading Libraries and Data

```
In [1]: import pandas as pd
import numpy as np
import seaborn as sns
from sqlalchemy import create_engine

import joblib
from sklearn.model_selection import train_test_split
from sklearn.model_selection import train_test_split
from sklearn.feature_extraction.text import TfidfVectorizer
import xgboost as xgb
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix
from imblearn.over_sampling import SMOTE
```

```
In [2]: host = r'127.0.0.1'
db = r'MSDS610'
user = r'postgres'
pw = r'postgres'
port = r'5432'
schema = r'clean'
```

```
In [3]: db_conn = create_engine("postgresql://{user}:{pw}@{host}:{port}/{db}".format(user, pw, host, port, db))
```

```
In [4]: table_name = r'movies_cleaned'
schema = r'cleaned'
```

```
In [5]: df = pd.read_sql_table(table_name, db_conn, schema)
```

```
In [6]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4803 entries, 0 to 4802
Data columns (total 16 columns):
#   Column                Non-Null Count  Dtype
---  -
0   budget                4803 non-null   int64
1   popularity            4803 non-null   float64
2   revenue               4803 non-null   int64
3   title                 4803 non-null   object
4   vote_average          4803 non-null   float64
5   vote_count            4803 non-null   int64
6   clean_genres           4803 non-null   object
7   clean_keywords         4803 non-null   object
8   clean_tagline          4803 non-null   object
9   clean_overview         4803 non-null   object
10  financial_success       4803 non-null   int64
11  critical_success        4803 non-null   int64
12  audience_success        4803 non-null   int64
13  budget_score            4803 non-null   float64
14  success_score           4803 non-null   float64
15  worth_funding           4803 non-null   int64
dtypes: float64(4), int64(7), object(5)
memory usage: 600.5+ KB
```

## Splitting the Data (3-way)

```
In [7]: text_features = ["clean_genres", "clean_keywords", "clean_tagline", "clean_overview"]
        y = df["worth_funding"]
```

```
In [8]: df["combined_text"] = df[text_features].apply(lambda x: " ".join(x.astype(str)), axis=1)
```

```
In [9]: X_train, X_temp, y_train, y_temp = train_test_split(df["combined_text"], y, test_size=0.2, random_state=4, stratify=y)
        X_val, X_test, y_val, y_test = train_test_split(X_temp, y_temp, test_size=0.5, random_state=4, stratify=y_temp)
```

```
In [10]: X_temp.info()
```

```
<class 'pandas.core.series.Series'>  
Index: 961 entries, 155 to 3088  
Series name: combined_text  
Non-Null Count  Dtype  
-----  
961 non-null    object  
dtypes: object(1)  
memory usage: 15.0+ KB
```

```
In [11]: X_temp.head()
```

```
Out[11]: 155    fantasy action adventure family comedy dyr...  
1953    comedy drama romance sex aquarium nudity ...  
1567    mystery thriller chicago fbi menace covere...  
3770    adventure action western robbery miner tre...  
3905    drama foreign female nudity pagan aztec ind...  
Name: combined_text, dtype: object
```

```
In [12]: y_temp.info()
```

```
<class 'pandas.core.series.Series'>  
Index: 961 entries, 155 to 3088  
Series name: worth_funding  
Non-Null Count  Dtype  
-----  
961 non-null    int64  
dtypes: int64(1)  
memory usage: 15.0 KB
```

```
In [13]: y_temp.sample(10)
```

```
Out[13]: 1452    2  
4099    2  
621     0  
3293    0  
2701    1  
480     1  
2333    0  
2306    1  
3738    1  
1169    0  
Name: worth_funding, dtype: int64
```

```
In [14]: y_temp.value_counts()
```

```
Out[14]: worth_funding
0      482
1      354
2      125
Name: count, dtype: int64
```

```
In [15]: print(X_test.shape)
X_test.head()
```

```
(481,)
Out[15]: 2860    drama thriller mystery toronto lesbian rem...
1843    drama family fantasy adventure based on nov...
366     comedy politics politician election campaign...
4272    drama mystery secret nightclub pet shop in...
2447    drama mystery thriller audio tape hitman p...
Name: combined_text, dtype: object
```

```
In [16]: print(X_val.shape)
X_val.head()
```

```
(480,)
Out[16]: 3932    action comedy horror monster pub duringcre...
3440    comedy drama mystery independent film every ...
1417    action comedy crime new york money launderi...
3133    comedy hotel infidelity onenight stand frie...
2397    comedy alcohol baby party family fraternit...
Name: combined_text, dtype: object
```

## Building the Model

```
In [17]: vectorizer = TfidfVectorizer(max_features=12000)
X_train_tfidf = vectorizer.fit_transform(X_train)
X_val_tfidf = vectorizer.transform(X_val)
```

```
In [18]: smote = SMOTE(sampling_strategy="auto", random_state=4)
X_train_resampled, y_train_resampled = smote.fit_resample(X_train_tfidf, y_train)
```

```
In [19]: xgb_model = xgb.XGBClassifier(  
        objective="multi:softprob",  
        eval_metric="mlogloss",  
        reg_lambda=6.0,  
        reg_alpha=4.0,  
        max_depth=3,  
        learning_rate=0.03,  
        n_estimators=500,  
        subsample=0.7,  
        colsample_bytree=0.7)  
xgb_model.fit(X_train_resampled, y_train_resampled)
```

```
Out[19]: ▼ XGBClassifier ⓘ  
XGBClassifier(base_score=None, booster=None, callbacks=None,  
              colsample_bylevel=None, colsample_bynode=None,  
              colsample_bytree=0.7, device=None, early_stopping_rounds=None,  
              enable_categorical=False, eval_metric='mlogloss',  
              feature_types=None, gamma=None, grow_policy=None,  
              importance_type=None, interaction_constraints=None,  
              learning_rate=0.03, max_bin=None, max_cat_threshold=None,  
              max_cat_to_onehot=None, max_delta_step=None, max_depth=3,  
              max_leaves=None, min_child_weight=None, missing=nan,  
              monotone_constraints=None, multi_strategy=None, n_estimators=500,  
              n_jobs=None, num_parallel_tree=None, objective='multi:softprob', ...)
```

```
In [20]: y_train_pred = xgb_model.predict(X_train_resampled)
```

```
In [21]: train_accuracy = accuracy_score(y_train_resampled, y_train_pred)  
train_report = classification_report(y_train_resampled, y_train_pred)  
train_conf_matrix = confusion_matrix(y_train_resampled, y_train_pred)
```

```
In [22]: print(f"Training Accuracy: {train_accuracy:.4f}")  
print("Training Classification Report:\n", train_report)  
print("Training Confusion Matrix:\n", train_conf_matrix)
```

Training Accuracy: 0.7229

Training Classification Report:

	precision	recall	f1-score	support
0	0.60	0.82	0.69	1926
1	0.77	0.57	0.65	1926
2	0.88	0.78	0.82	1926
accuracy			0.72	5778
macro avg	0.75	0.72	0.72	5778
weighted avg	0.75	0.72	0.72	5778

Training Confusion Matrix:

```
[[1588 292  46]
 [ 673 1092 161]
 [ 387  42 1497]]
```

## Exporting Data for Validation (Notebook 2)

### Saving the Model and the Vectorizer

```
In [23]: joblib.dump(xgb_model, "xgb_model.pkl")
         joblib.dump(vectorizer, "tfidf_vectorizer.pkl")
```

```
Out[23]: ['tfidf_vectorizer.pkl']
```

### Saving the Validation Set

```
In [24]: X_val.to_csv("X_val.csv", index=False, header=True)
         y_val.to_csv("y_val.csv", index=False, header=True)
```

```
In [39]: df.to_csv("clean_movie_df.csv", index=False)
```